

OpenACC

OpenMP®



*Related Programming
Models*

07.06.2023




- CUDA programming model
 - C / C++
 - Fortran
- CUDA API
 - Runtime
 - Driver
- CUDA parallel computing platform
 - Architectures: Tesla, Fermi, Kepler, Maxwell, Pascal, Volta, Turing, Ampere, ...
 - PTX ISA

NVIDIA only




- Standards by the Khronos Group
 - OpenCL: cross-platform, parallel programming
 - SYCL: higher-level programming model for OpenCL
 - Vulkan: low-level graphics and compute shader API
- Standard by the HSA foundation
 - HSA (Heterogeneous System Architecture)
- Directive-based standards
 - OpenACC
 - OpenMP 4.0
- Vendor specific
 - C++ AMP (Accelerated Massive Parallelism, Microsoft) – deprecated
 - DPC++ (Data Parallel C++, Intel)
 - ROCm (Radeon Open Compute, AMD)
 - CUDA (Compute Unified Device Architecture, NVIDIA)



- Initiated by Apple (2008)
- Maintained by the Khronos Group
- Goal:
 - Cross-platform: desktop, embedded, HPC
 - Cross-device: CPU, GPU, FPGA, DSP, ...
- Problem:
 - Slow adoption of new features
 - Functional, but no performance portability

OpenCL Conformant Implementations



Vendor timelines are first conformant submission for each spec generation

Dec08
OpenCL 1.0
Specification

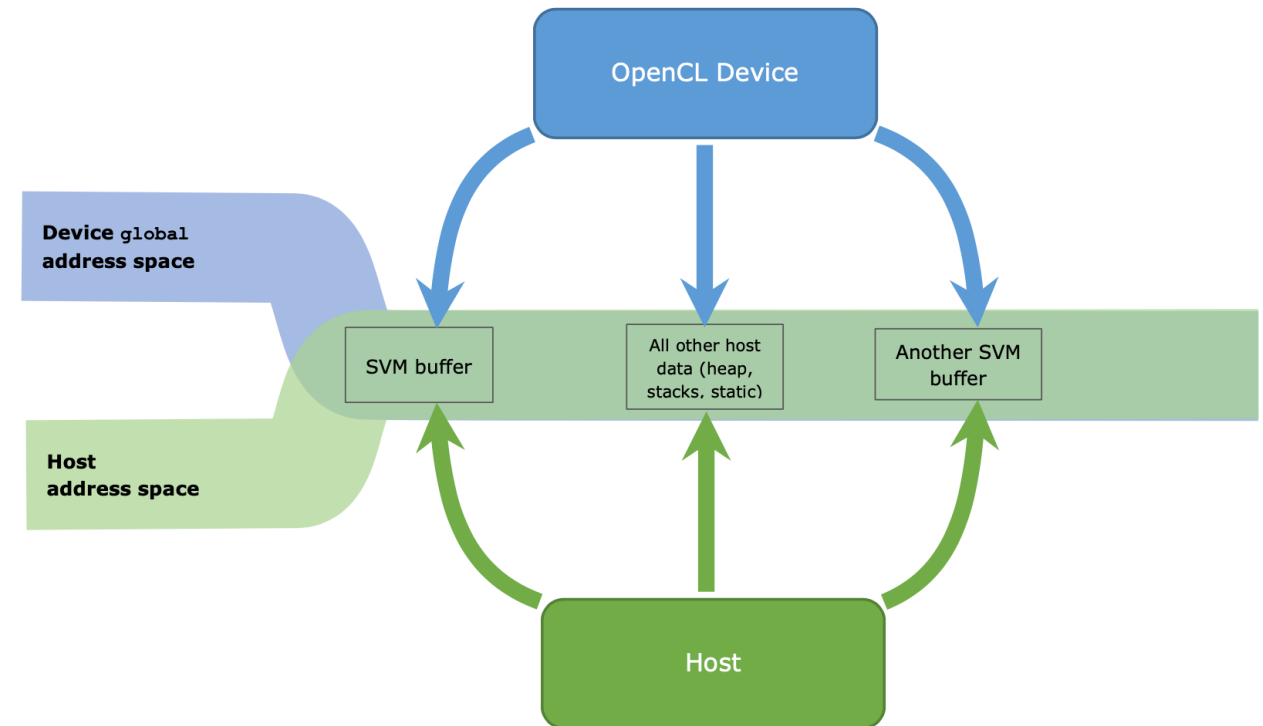
Jun10
OpenCL 1.1
Specification

Nov11
OpenCL 1.2
Specification

Nov13
OpenCL 2.0
Specification

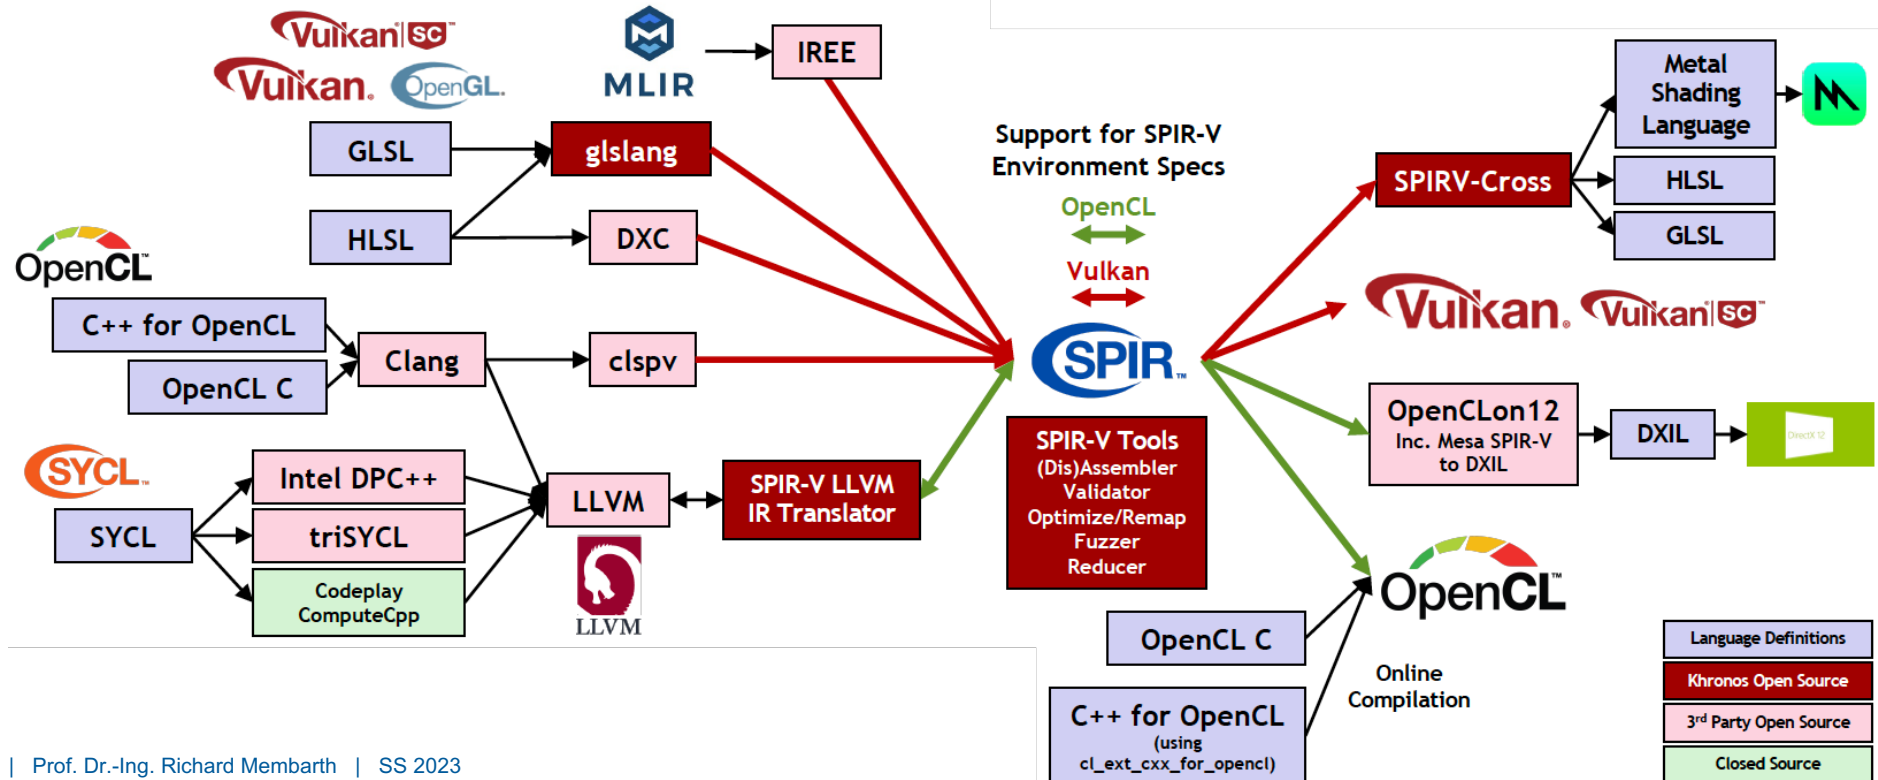
Nov15
OpenCL 2.1
Specification

- OpenCL 1.2 (11/2011)
 - C99 based kernel language
 - Most prevalent standard
- OpenCL 2.0 (11/2013)
 - Shared virtual memory (SVM)
 - Coarse grained (required)
 - Fine grained (optional)
 - System (optional)
 - C11 atomics
 - Ordering/consistency: relaxed, acquire, release, acq_rel, seq_cst
 - Memory scope: work_item, work_group, device, all_svm_devices






- OpenCL 2.1 (11/2015)
 - C++14 based kernel language (optional)
 - SPIR-V intermediate representation



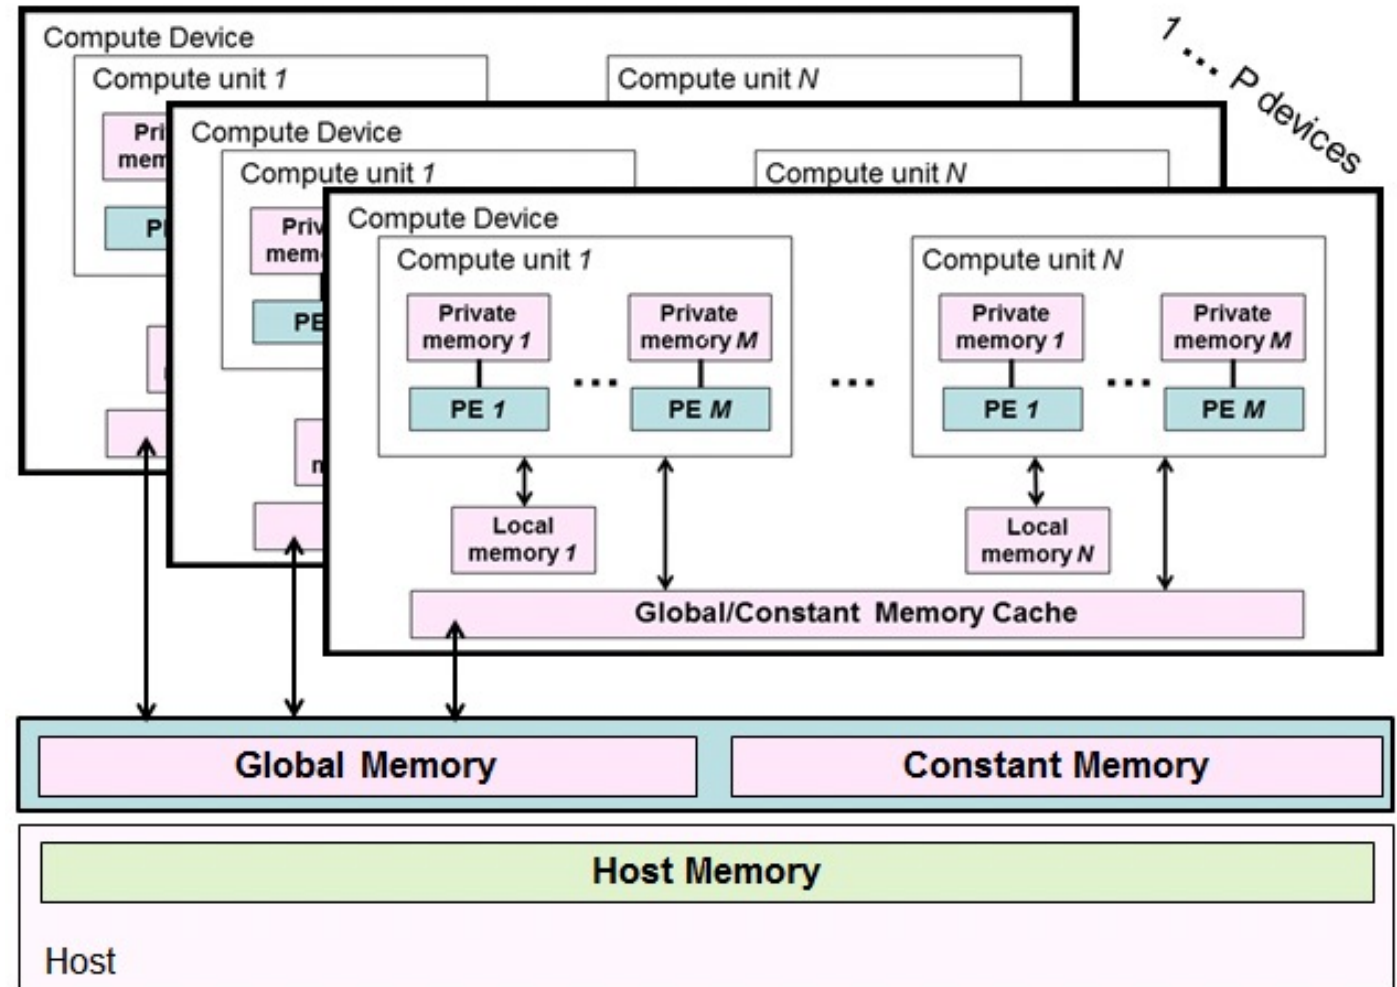


- OpenCL 2.2 (05/2017)
 - C++14 based kernel language (core)
- OpenCL 3.0 (09/2020)
 - OpenCL 1.2 mandatory
 - OpenCL 2.x features as optional modules
- Next: launch OpenCL kernels via Vulkan ?

- Single host system
- Multiple compute devices
 - CPU, GPU, FPGA, ...
- Compute unit → CUDA multiprocessor
- Processing element → CUDA core



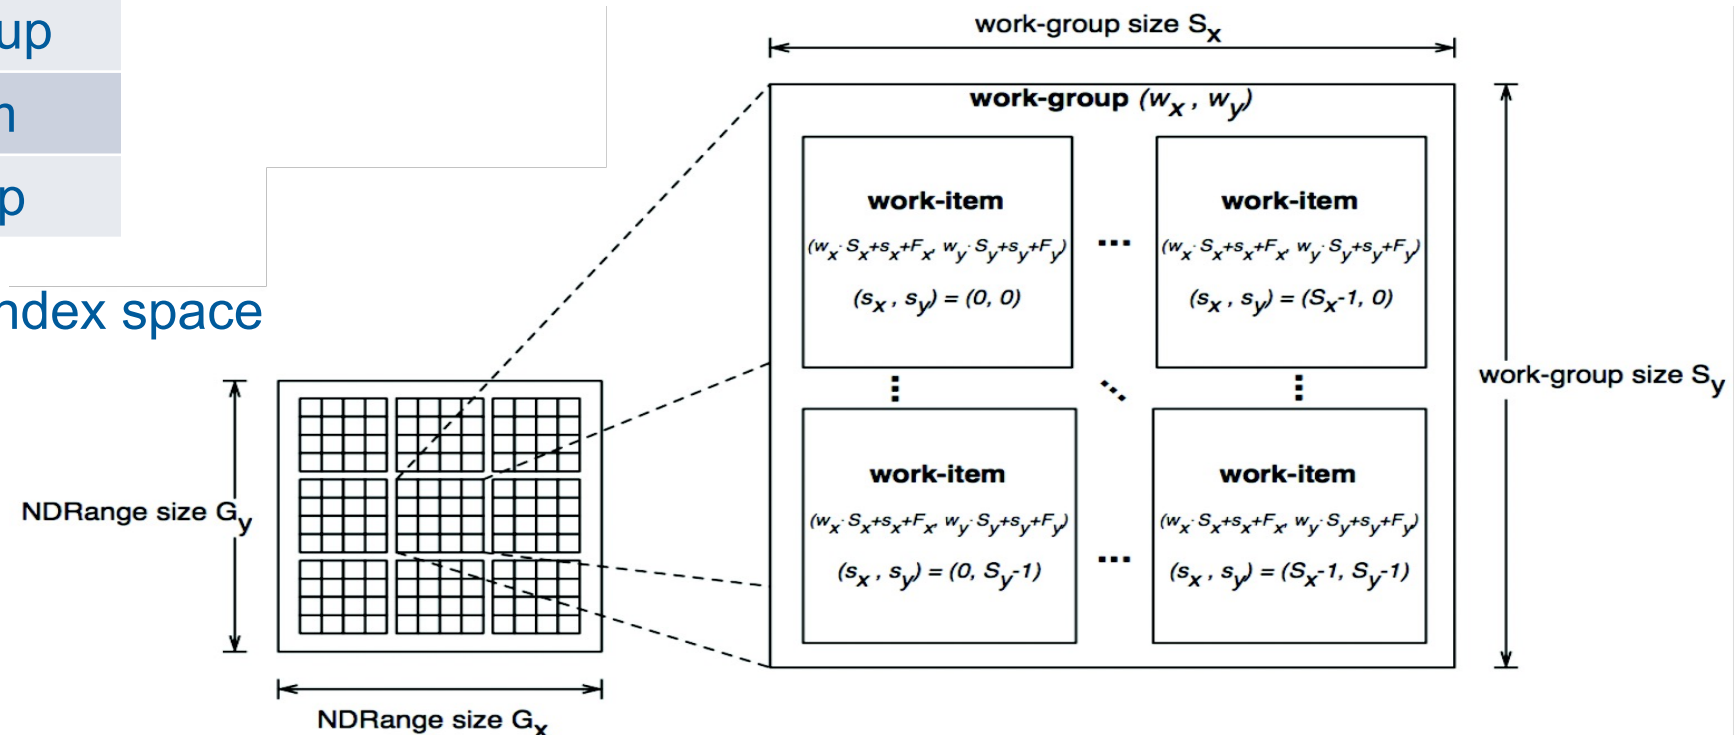
- Similar to CUDA
 - Host memory
 - Global memory
 - Constant memory
 - Local memory
 - Private memory



- Index space mapping similar to CUDA

CUDA	OpenCL
grid	NDRange
block	work-group
thread	work-item
warp	sub-group

- Kernel launch specifies index space



- Kernel language similar to CUDA

	CUDA	OpenCL
thread index	threadIdx.x	get_local_id(0)
block index	blockIdx.x	get_group_id(0)
block dimension	blockDim.x	get_local_size(0)
grid dimension	gridDim.x	get_global_size(0)
global memory	__global__	__global
group memory	__shared__	__local
constant memory	__constant__	__constant
barrier	__syncthreads	barrier(CLK_LOCAL_MEMFENCE)



- Generated code often incorrect for complex code
- Deprecated since macOS Mojave (10.14)
- Canceled for macOS Catalina (10.15) via security update 2021-002

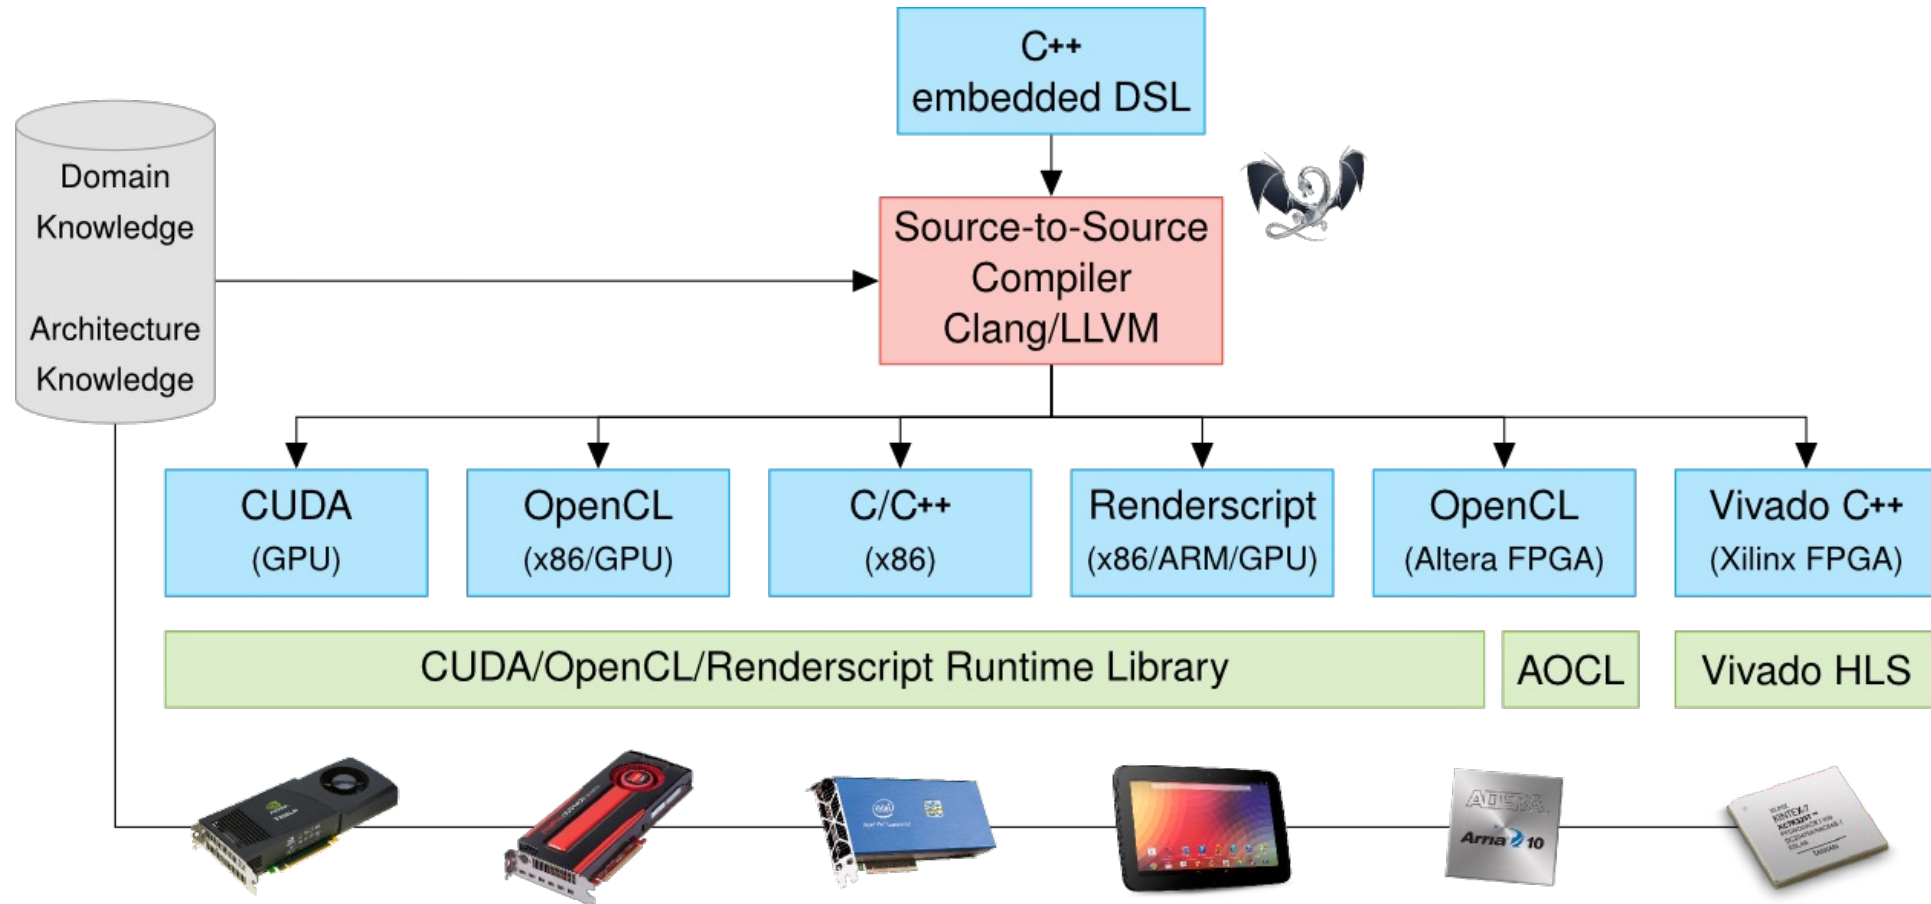
- Metal
 - Graphics and compute shader API
 - Replaces OpenGL and OpenCL



- Not available on Android
- Android provides RenderScript / FilterScript
 - Compute API, C99-derived language
 - Java class reflected by compiler
 - Mainly in maintenance mode since years
 - Introduced with Android 3 (2011)
 - Deprecated in Android 12 (2021)
- Vulkan
 - Supported since Android 7 (2016)
 - Replaces RenderScript / FilterScript

android

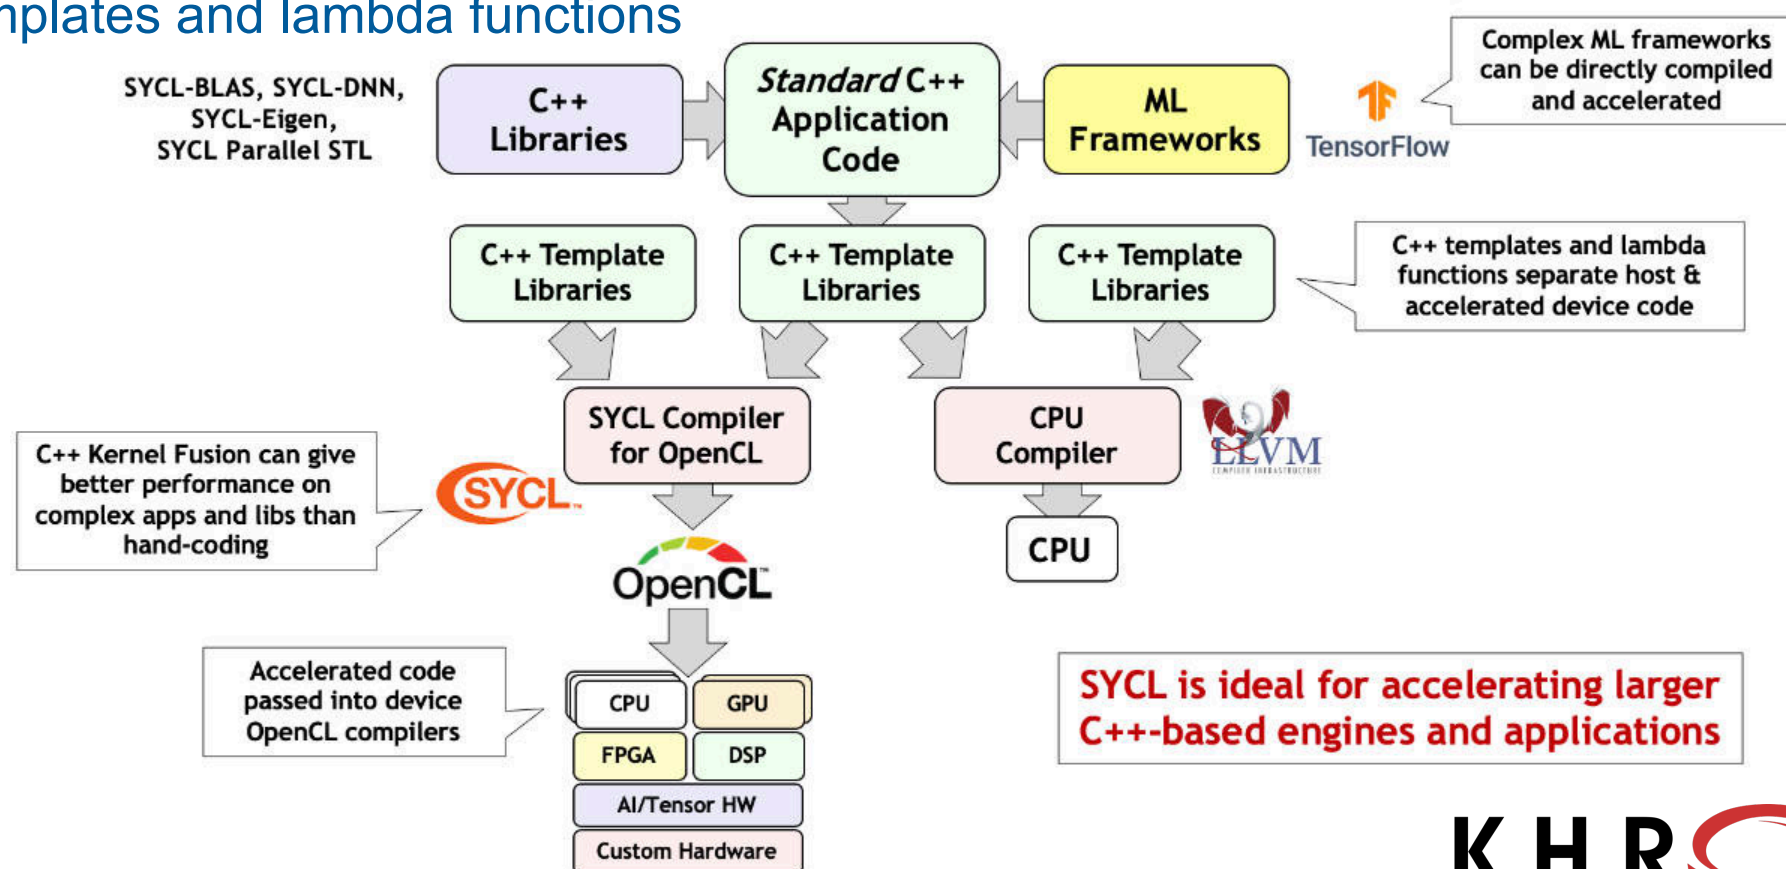




<https://hipacc-lang.org>

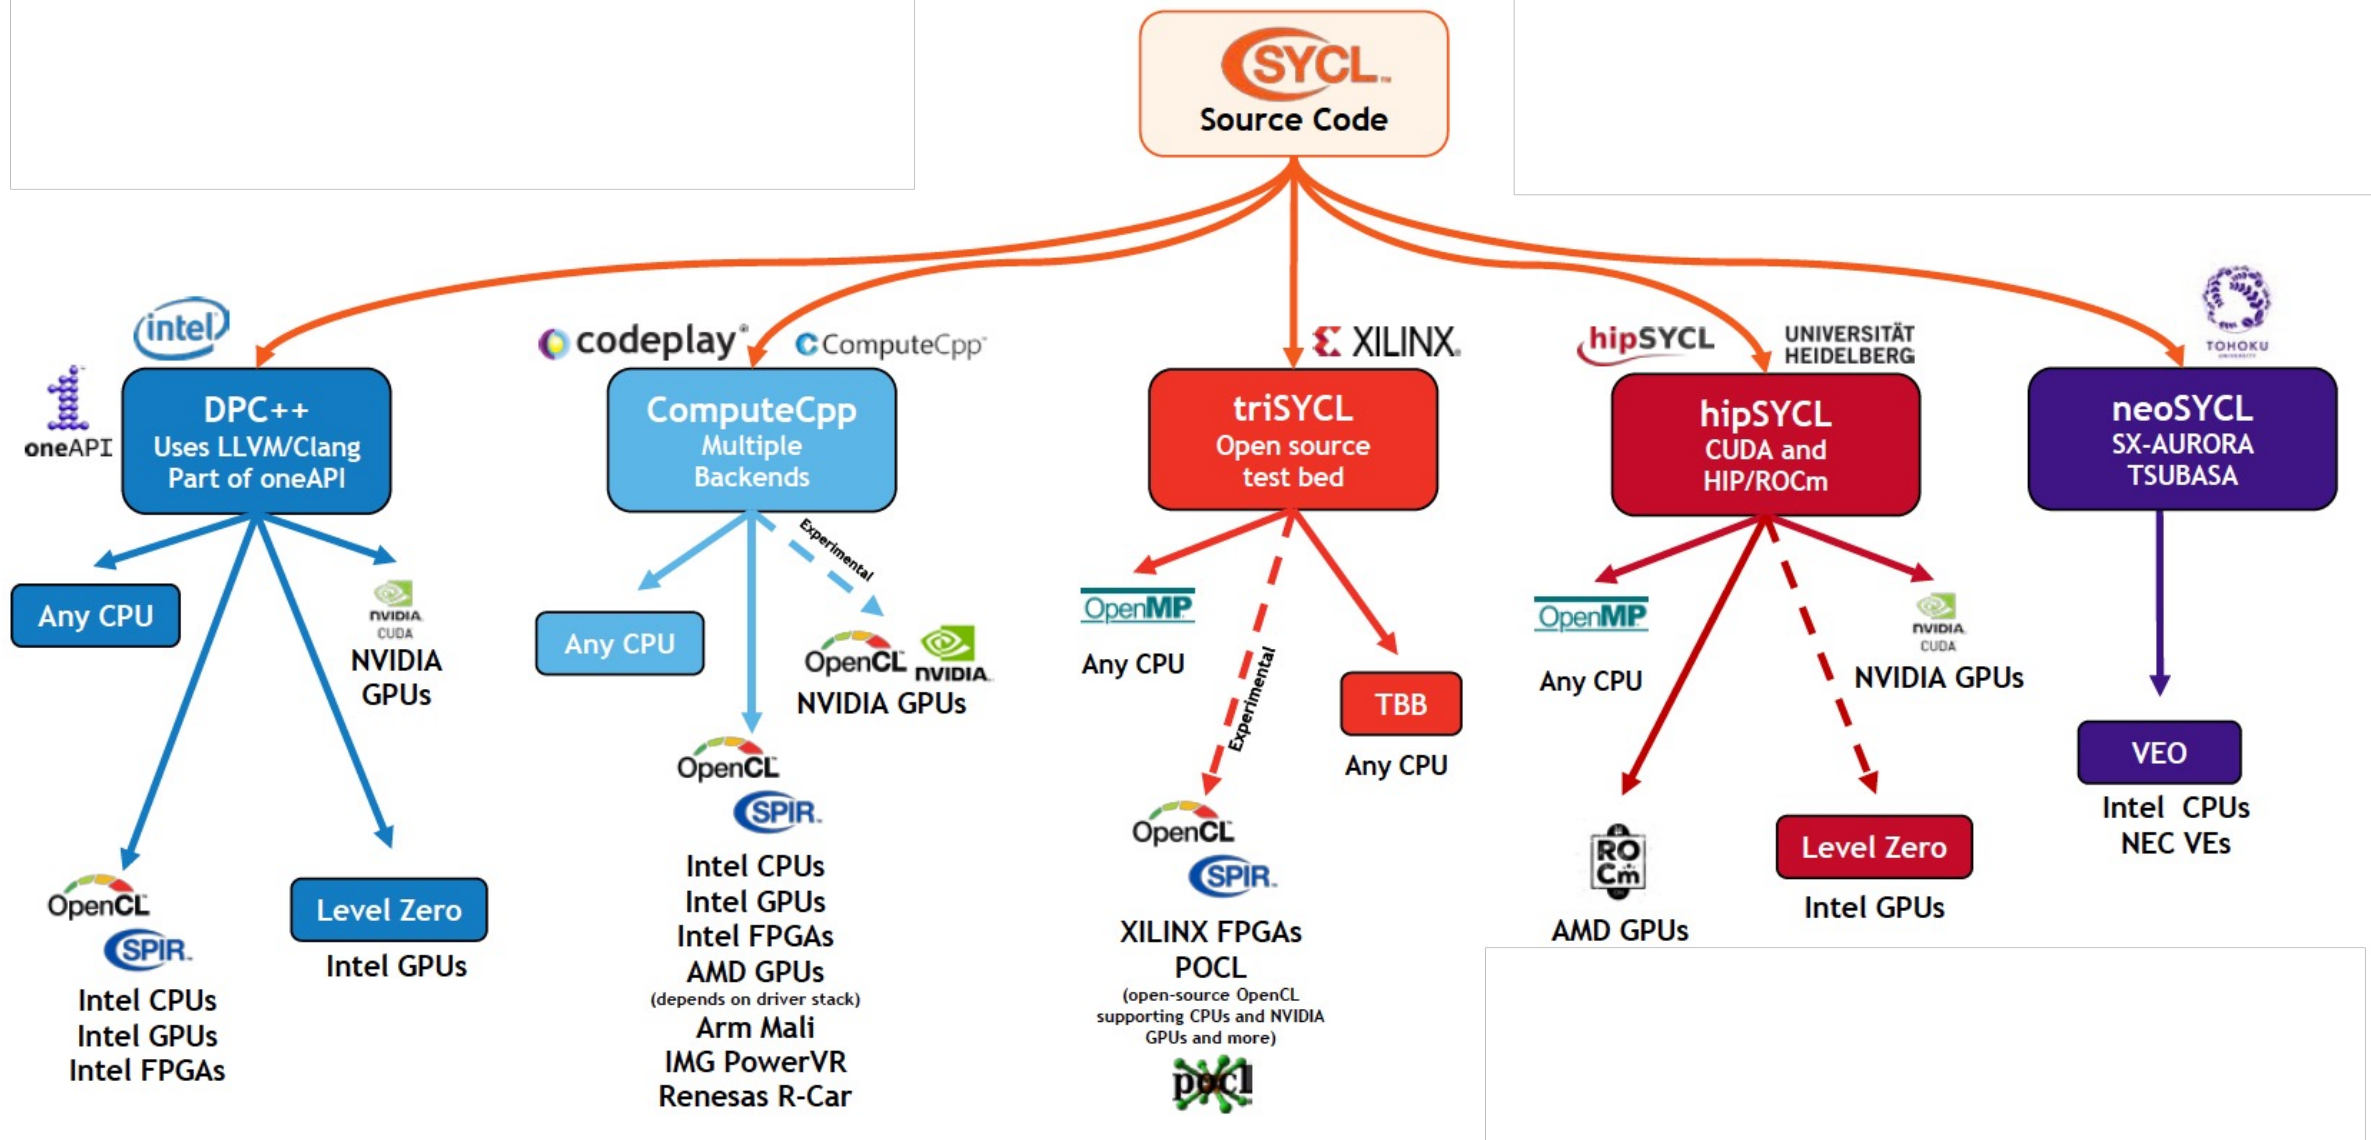


- C++ single source heterogeneous programming
- Templates and lambda functions



SYCL

Backends



```

#include <SYCL/sycl.hpp>
class vec_add;

int main() {
    auto array_size = 256;
    std::vector<float> A(array_size, 1.0f);
    std::vector<float> B(array_size, 1.0f);
    std::vector<float> C(array_size);
    { // beginning C++ scope - SYCL takes ownership of memory
        auto sycl_queue = cl::sycl::queue;
        auto A_buff = cl::sycl::buffer<float>(A.data(), cl::sycl::range<1>(array_size));
        auto B_buff = cl::sycl::buffer<float>(B.data(), cl::sycl::range<1>(array_size));
        auto C_buff = cl::sycl::buffer<float>(C.data(), cl::sycl::range<1>(array_size));
        auto num_groups = sycl_queue.get_device().get_info<cl::sycl::info::device::max_compute_units>();
        auto work_group_size = sycl_queue.get_device().get_info<cl::sycl::info::device::max_work_group_size>();
        auto total_threads = num_groups * work_group_size;
        sycl_queue.submit([&](cl::sycl::handler &cgh) {
            auto A_acc = A_buff.get_access<cl::sycl::access::mode::read>(cgh);
            auto B_acc = B_buff.get_access<cl::sycl::access::mode::read>(cgh);
            auto C_acc = C_buff.get_access<cl::sycl::access::mode::write>(cgh);
            cgh.parallel_for<class vec_add>(
                cl::sycl::range<1>{total_threads}, [=](cl::sycl::item<1> index) {
                    C_acc[index] = A_acc[index] + B_acc[index];
                }
            );
        });
    } // end of C++ scope
}

```



- Heterogeneous System Architecture (HSA)
- *Optimized platform architecture for OpenCL*
 - Coherent memory model
 - HSAIL intermediate representation
 - Fully asynchronous dispatch and runtime
- Only publicly available implementation
 - AMD ROCm



- Open-Source platform for HPC
- Collection of tools, libraries, frameworks
 - HCC: single source C++ accelerator language (deprecated)
 - HIP: runtime / kernel language for AMD/NVIDIA GPUs
 - cudaMalloc → hipMalloc
 - HIPIFY → convert CUDA to HIP
 - OpenCL, OpenMP, ...
- Runtime implements HSA standard





- Early approaches (~2008)
 - Portland Group: PGI Accelerator
 - CAPS: HMPP Workbench
- Defined new standard (~2011)
 - OpenACC
- OpenMP 4.0 (~2013)
 - Support for offloading

OpenACC

OpenMP[®]



■ Pragma applies to next loop only

```
#pragma acc parallel loop
for (int i=0; i<length; ++i) {
    // sequential
    for (int j=0; j<length; ++j) {
        ...
    }
}
```

```
#pragma acc parallel loop
for (int i=0; i<length; ++i) {
    #pragma acc loop
    for (int j=0; j<length; ++j) {
        ...
    }
}
```

■ Region defines scope

```
#pragma acc region
{
    #pragma acc parallel loop
    for (int i=0; i<length; ++i) {
        #pragma acc loop
        for (int j=0; j<length; ++j) {
            ...
        }
    }
}
```



- Data regions can be specified for region
 - copy, copyin, copyout, create, present

```
#pragma acc data copy(x[0:N])
#pragma acc parallel loop
{
  for (int i=0; i<length; ++i) {
    #pragma acc loop
    for (int j=0; j<length; ++j) {
      ...
    }
  }
  ...
}
```

```
#pragma acc parallel loop copy(x[0:N])
for (int i=0; i<length; ++i) {
  #pragma acc loop
  for (int j=0; j<length; ++j) {
    ...
  }
}
```



- OpenACC provides set of
 - Compiler directives (pragmas)
 - Library routines (runtime)
 - Environment variables

```
void vec_add(float* __restrict__ out, const float* in1, const float* in2, int lenght) {  
    #pragma acc parallel loop copyin(in1[0:length], in2[0:length]), copyout(out[0:length])  
    for (int i=0; i<length; ++i) {  
        out[i] = in1[i] + in2[i];  
    }  
}
```




- WebCL
 - JavaScript bindings to OpenGL
 - Provided by the Khronos Group
 - No native support by browsers
- WebGPU
 - Future web standard and JavaScript API
 - APIs for accelerated graphics and compute
 - Initiated by Apple
 - Support by W3C, Mozilla, Microsoft, Google, ...