



17. APRIL 2023

## GPU PROGRAMMING ASSIGNMENT 3

Submission deadline for the exercises: 24. April 2023

### 3.1 Memory Spaces

- a) The following table lists all memory spaces available in CUDA as well properties such as
- scope (is it private to each thread, can all threads within a block access it, or is it globally accessible?);
  - whether it is read-only in device code;
  - whether it is cached;
  - can it be dynamically allocated from host or device?

Unfortunately, the table is incomplete. Please fill in the missing information.

name	scope (thread/block/global)	read-only (y/n)	cached (y/n)	dyn. alloc. from	
				host (y/n)	device (y/n)
	global	n	y	y	y
constant					
texture					
	block				
register					
	thread				
system					

### 3.2 Gaussian Blur Filter

The purpose of this exercise is to get familiar with some of the different memory spaces in CUDA using a Gaussian blur filter on random data.

- a) The `blur.cu` file contains the host and device code for the Gaussian blur filter. Implement the `blur_kernel` CUDA kernel that applies a given filter `filter` to an input image `in` and stores the result to `out`:

```

1 __global__ void blur_kernel(float* in, float* out, float* filter,
2                             int size_x, int size_y, int width, int height) {
3     // ...
4 }

```

The filter is of size  $size\_x \times size\_y$  and the images have  $width \times height$  pixels. We use a simplified boundary handling as shown in the `compute_reference` implementation: When we have an out-of-bounds memory access, we skip the computation for that filter coefficient. Add the missing CUDA API calls on the host side in order to launch the kernel:

- allocate device memory
- copy host memory to the device
- launch the kernel
- copy the device memory back to the host
- free device memory

Make sure that your code checks for CUDA API errors and works for different kernel input sizes.

- b) Benchmark the Gaussian blur filter reporting the average execution time of 10 kernel invocations using either the CUDA event API or `nsys profile --stats=true` exploring different memory spaces:

- adding the `const` qualifier for kernel parameters where appropriate
- adding the `__restrict__` qualifier for kernel parameters where appropriate
- adding the `const` and `__restrict__` qualifier for kernel parameters where appropriate
- using `__constant__` memory for filter
- using texture memory for in

For the texture memory, define your texture using the `cudaReadModeElementType`:

```

1 texture<float, 2, cudaReadModeElementType> in_tex;

```

Allocate a CUDA array for texture and bind the CUDA array to the texture:

```

1 cudaChannelFormatDesc desc = cudaCreateChannelDesc<float>();
2 cudaArray* d_arr_in;
3 cudaMallocArray(&d_arr_in, &desc, width, height);
4 cudaMemcpy2DToArray(d_arr_in, ...);
5 cudaBindTextureToArray(&in_tex, d_arr_in, &desc);

```

For further documentation, lookup the functions of the CUDA API on <https://docs.nvidia.com/cuda/cuda-runtime-api/index.html>.